

Operating System Concepts (cont'd)

Course Code: CSC 2209

Course Title: Operating Systems



Dept. of Computer Science
Faculty of Science and Technology

Lecturer No:	02	Week No:	02	Semester:	Fall 2024-25
Assistant Professor	<i>Dr. Rajarshi Roy Chowdhury, rajarshi@aiub.edu</i>				

Lecture Outline



1. I/O Structure
2. Storage Structure
3. How a Modern Computer Works
4. Direct Memory Access Structure
5. Computer-System Architecture
6. A Dual-Core Design
7. Clustered Systems
8. Operating-System Operations
9. Multiprogramming and Multitasking
10. Dual-mode and Multimode Operation
11. Process Management
12. Memory Management
13. File-system Management
14. Caching

I/O Structure

- ❑ After I/O starts, control returns to user program only upon I/O completion
 - ❑ **Wait instruction** idles the CPU until the next interrupt
 - ❑ Wait loop (contention for memory access)
 - ❑ At most one I/O request is outstanding at a time, no simultaneous I/O processing
- ❑ After I/O starts, control returns to user program **without waiting for I/O completion**
 - ❑ **System call** – request to the OS to allow user to wait for I/O completion
 - ❑ **Device-status table** contains entry for each I/O device indicating its **type, address, and state**
 - ❑ OS indexes into **I/O device table** to determine device status and to modify table entry to include interrupt


I/O Structure

Device-status table

Device ID	Device Type	Address	State
0	Disk Drive	0x1000	Busy
1	Printer	0x1001	Idle
2	Keyboard	0x1002	Ready
3	Network Card	0x1003	Error

Storage Structure

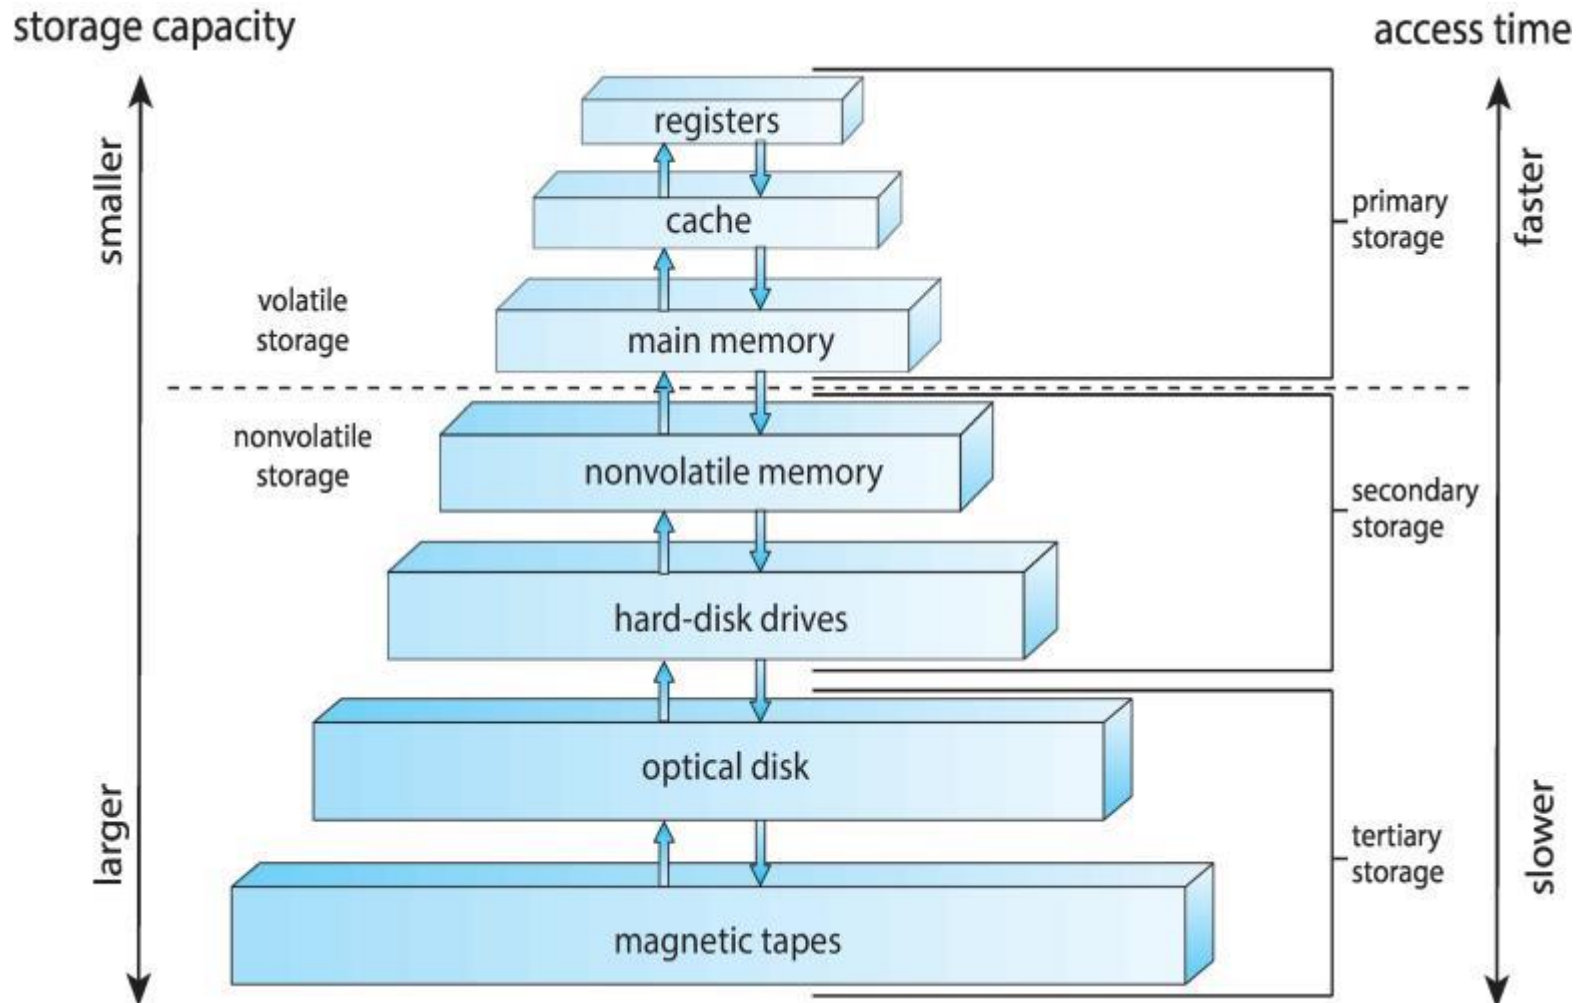
- ❑ **Main memory** – only large storage media that the CPU can access directly
 - ❑ **Random access**
 - ❑ Typically **volatile**
 - ❑ Typically **random-access memory (RAM)** in the form of **Dynamic Random-access Memory (DRAM)**
- ❑ **Secondary storage** – extension of main memory that provides large **nonvolatile** storage capacity
- ❑ **Hard Disk Drives (HDD)** – **rigid metal or glass platters** covered with magnetic recording material
 - ❑ **Disk surface** is **logically** divided into **tracks**, which are subdivided into **sectors**
 - ❑ The **disk controller** determines the **logical interaction between the device and the computer**
- ❑ **Non-volatile memory (NVM)** devices– faster than hard disks,
 - ❑ Various technologies
 - ❑ Becoming more popular as capacity and performance increases, price drops



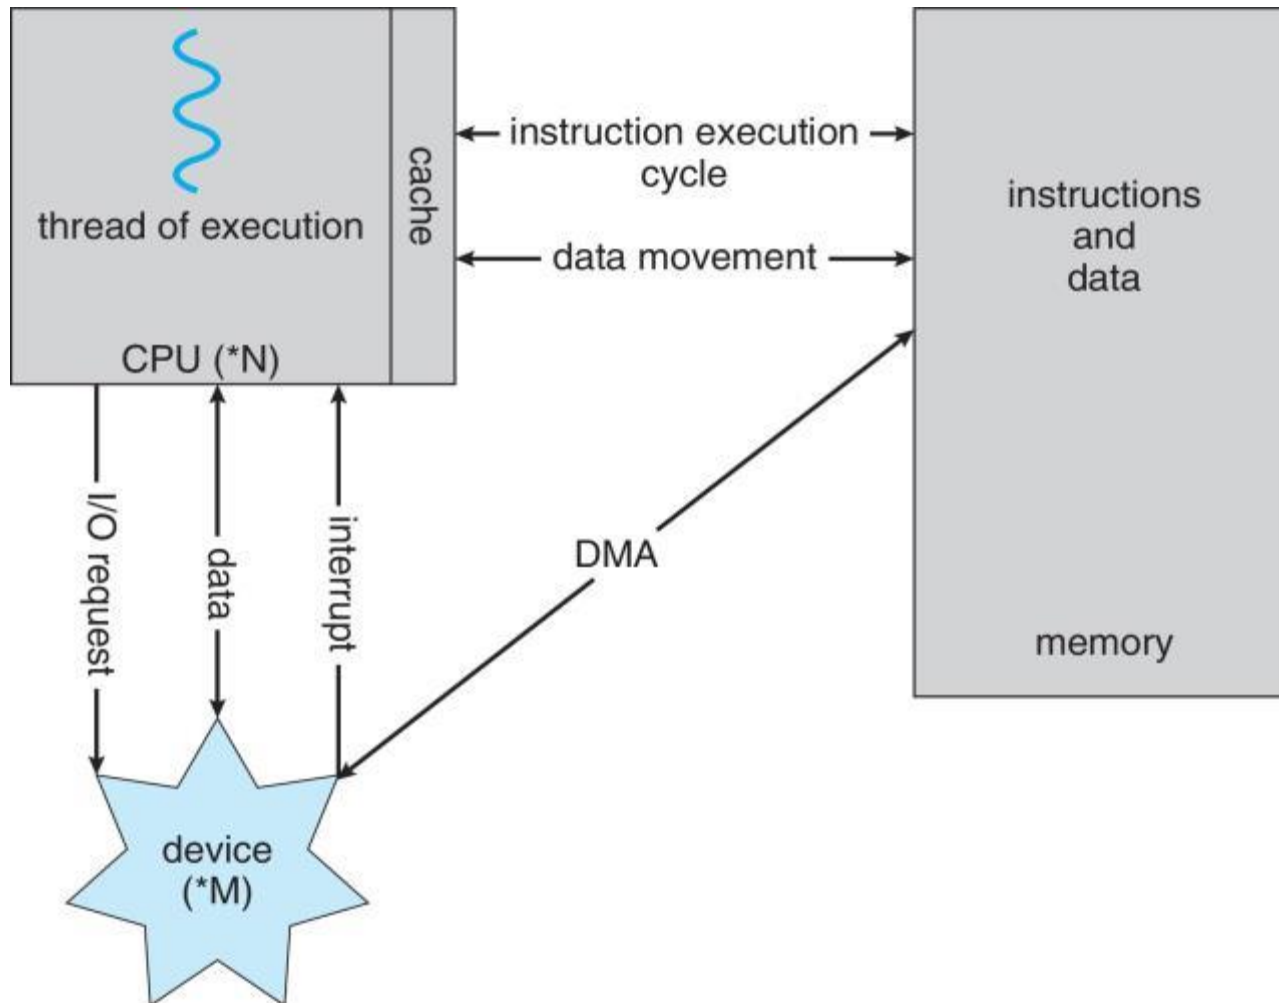
Storage Hierarchy

- ❑ Storage systems organized in hierarchy
 - ❑ Speed
 - ❑ Cost
 - ❑ Volatility
- ❑ **Caching** – copying information into faster storage system; **main memory can be viewed as a cache for secondary storage**
- ❑ **Device Driver** for each device controller to manage I/O
 - ❑ Provides uniform **interface between controller and kernel**

Storage-Device Hierarchy



How a Modern Computer Works



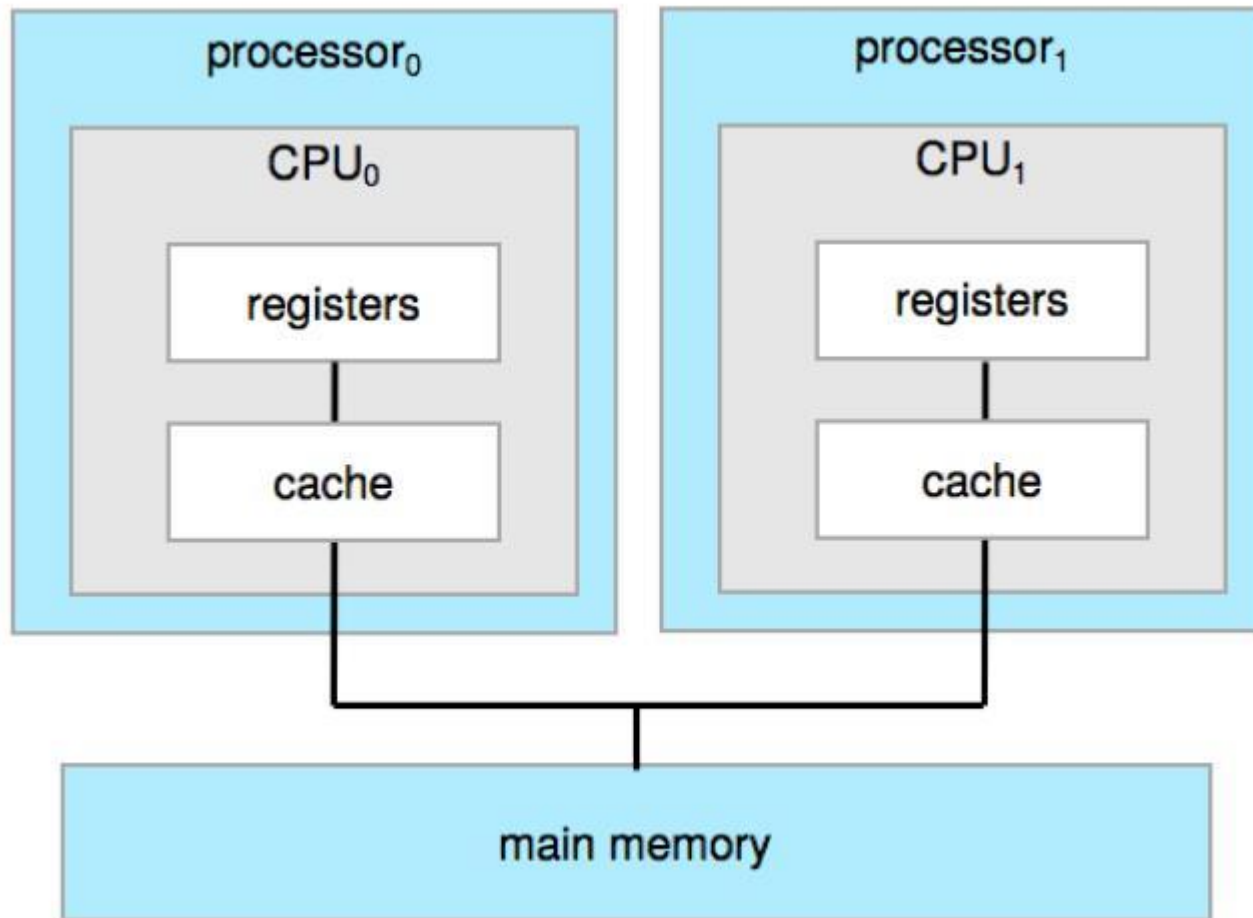
Direct Memory Access (DMA) Structure

- ❑ Used for **high-speed I/O devices** able to **transmit information at close to memory speeds**
- ❑ Device controller transfers **blocks of data** from buffer storage **directly to main memory without CPU intervention**
- ❑ **Only one interrupt is generated per block, rather than the one interrupt per byte**

Computer-System Architecture

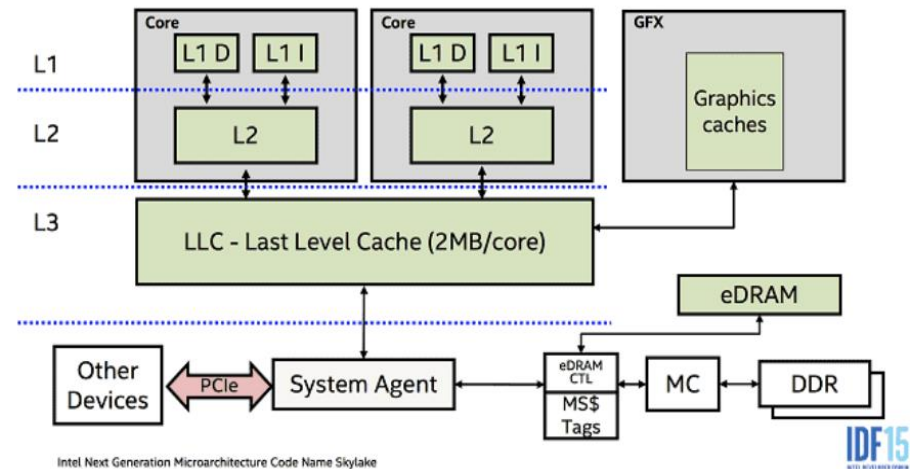

- ❑ Most systems use a single general-purpose processor
 - ❑ Most systems have special-purpose processors as well
- ❑ **Multiprocessors** systems growing in use and importance
 - ❑ Also known as **parallel systems**, **tightly-coupled systems**
 - ❑ Advantages include:
 1. **Increased throughput**
 2. **Economy of scale**
 3. **Increased reliability** – graceful degradation or fault tolerance
 - ❑ Two types:
 1. **Asymmetric Multiprocessing** – each processor is assigned a specific task.
 2. **Symmetric Multiprocessing** – each processor performs all tasks

Symmetric Multiprocessing Architecture



A Dual-Core Design

- ❑ Multi-chip and **multicore**
- ❑ Systems containing all chips
 - ❑ Chassis containing multiple separate systems




Cache Level	Speed	Size	Volatility	Purpose
L1 Cache	Fastest	16 KB - 128 KB	Volatile	Immediate access to data and instructions
L2 Cache	Fast	256 KB - 2 MB	Volatile	Buffer between L1 and L3/main memory
L3 Cache	Moderate	2 MB - 32 MB+	Volatile	Shared cache for multi-core CPUs

Clustered Systems

- ❑ Like multiprocessor systems, but **multiple systems working together**
 - ❑ Usually sharing storage via a **storage-area network (SAN)**
 - ❑ Provides a **high-availability** service which survives failures
 - ❑ **Asymmetric clustering** has one machine in hot-standby mode
 - ❑ **Symmetric clustering** has multiple nodes running applications, monitoring each other
 - ❑ Some clusters are for **high-performance computing (HPC)**
 - ❑ Applications must be written to use **parallelization**
 - ❑ Some have **distributed lock manager (DLM)** to avoid conflicting operations

Clustered Systems




Operating-System Operations

- ❑ **Bootstrap program** – simple code to initialize the system, load the kernel
- ❑ Kernel loads
- ❑ Starts **system daemons** (services provided outside of the kernel)
- ❑ Kernel **interrupt driven** (hardware and software)
 - ❑ Hardware interrupt by one of the devices
 - ❑ Software interrupt (**exception** or **trap**):
 - ❑ Software error (e.g., division by zero)
 - ❑ Request for operating system service – **system call**
 - ❑ Other process problems include infinite loop, processes modifying each other or the operating system

Multiprogramming and Multitasking

- ❑ **Multiprogramming (Batch system)** needed for efficiency
 - ❑ Single user cannot keep CPU and I/O devices busy at all times
 - ❑ Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - ❑ A subset of total jobs in system is kept in memory
 - ❑ One job selected and run via **job scheduling**
 - ❑ When it has to wait (for I/O for example), OS switches to another job
- ❑ **Timesharing (multitasking)** is **logical extension** in which **CPU switches jobs so frequently that users can interact with each job** while it is running, creating **interactive** computing
 - ❑ **Response time** should be < 1 second
 - ❑ Each user has at least one program executing in memory **process**
 - ❑ If several jobs ready to run at the same time **CPU scheduling**
 - ❑ If processes doesn't fit in memory, **swapping** moves them in and out to run
 - ❑ **Virtual memory** allows execution of processes not completely in memory

Memory Layout for Multiprogrammed System



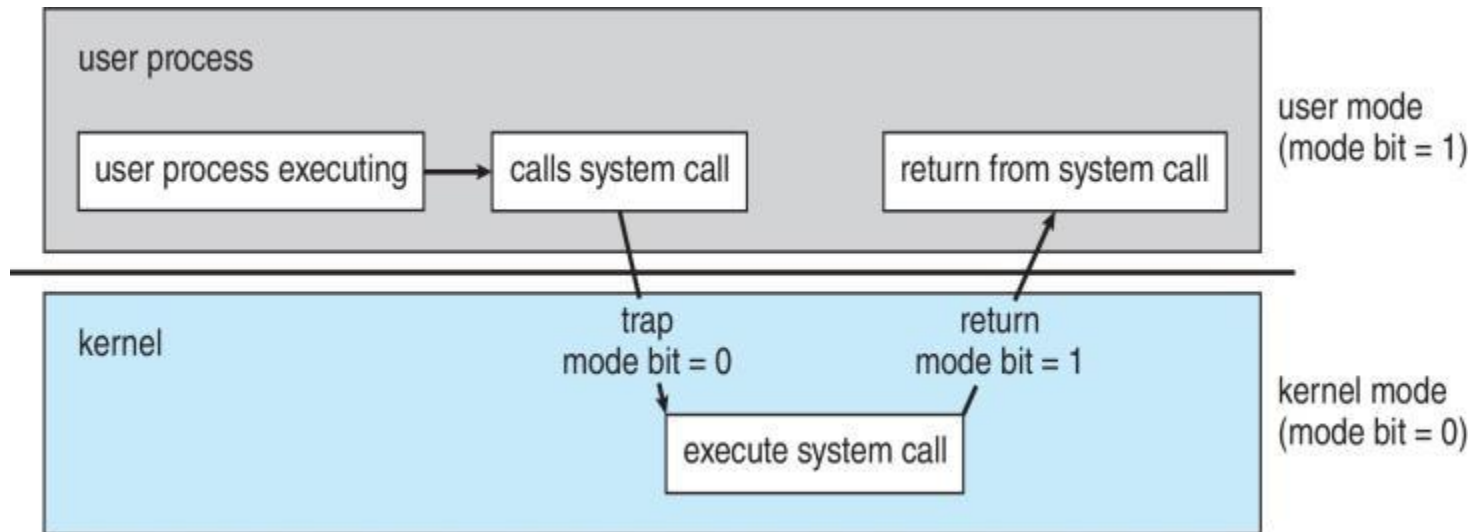
Dual-mode and Multimode Operation

- ❑ **Dual-mode** operation allows OS to protect itself and other system components
 - ❑ **User mode** and **kernel mode**
 - ❑ **Mode bit** provided by hardware
 - ❑ Provides ability to distinguish when system is running user code or kernel code
 - ❑ Some instructions designated as **privileged**, only executable in kernel mode
 - ❑ System call changes mode to kernel, return from call resets it to user
- ❑ Increasingly CPUs support multi-mode operations
 - ❑ i.e. **virtual machine manager (VMM)** mode for guest **VMs**

Transition from User to Kernel Mode

- ❑ Timer to prevent infinite loop / process hogging resources
 - ❑ Timer is set to interrupt the computer after some time period
 - ❑ Keep a counter that is decremented by the physical clock
 - ❑ Operating system set the counter (privileged instruction)
 - ❑ When counter zero generate an interrupt
 - ❑ Set up before scheduling process to regain control or terminate program that exceeds allotted time

Transition from User to Kernel Mode



Process Management

- ❑ A **process is a program in execution**. It is a **unit of work** within the system. Program is a *passive entity*, process is an *active entity*.
- ❑ Process needs resources to accomplish its task
 - ❑ CPU, memory, I/O, files
 - ❑ Initialization data
- ❑ Process termination requires reclaim of any reusable resources
- ❑ **Single-threaded process** has one **program counter** specifying location of next instruction to execute
 - ❑ Process **executes instructions sequentially, one at a time**, until completion
- ❑ **Multi-threaded process** has **one program counter per thread**
- ❑ Typically system has many processes, some user, some operating system running concurrently on one or more CPUs
 - ❑ Concurrency by multiplexing the CPUs among the processes / threads

Process Management Activities

The operating system is responsible for the following activities in connection with process management:

- ❑ **Creating and deleting** both user and system processes
- ❑ **Suspending and resuming** processes
- ❑ Providing mechanisms for **process synchronization**
- ❑ Providing mechanisms for **process communication**
- ❑ Providing mechanisms for **deadlock handling**

Memory Management

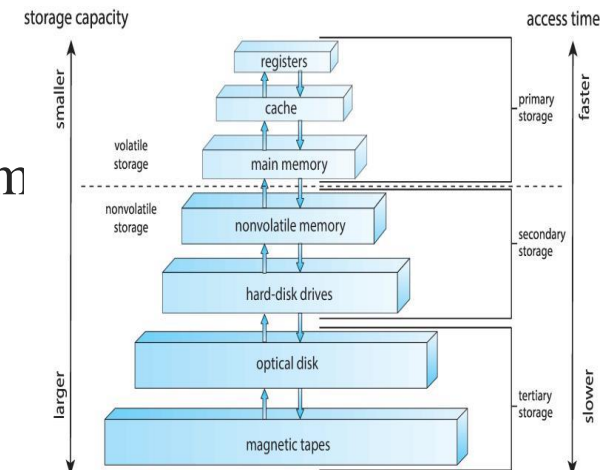
- ❑ To execute a program all (or part) of the **instructions must be in memory**
- ❑ All (or part) of the **data** that is **needed by the program must be in memory**
- ❑ **Memory management determines what is in memory and when**
 - ❑ Optimizing CPU utilization and computer response to users
- ❑ Memory management activities
 - ❑ **Keeping track** of which parts of memory are currently being used and by whom
 - ❑ Deciding which processes (or parts thereof) and data to move into and out of memory
 - ❑ **Allocating and deallocating** memory space as needed

File-system Management

- ❑ OS provides uniform, logical view of information storage
 - ❑ Abstracts physical properties to logical storage unit - **file**
 - ❑ Each medium is controlled by device (i.e., disk drive, tape drive)
 - ❑ Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)
- ❑ File-System management
 - ❑ Files usually **organized into directories**
 - ❑ **Access control** on most systems to **determine who can access** what
 - ❑ OS activities include
 - ❑ **Creating and deleting** files and directories
 - ❑ Primitives to **manipulate** files and directories
 - ❑ **Mapping files** onto secondary storage
 - ❑ **Backup files** onto stable (non-volatile) storage media

Caching

- ❑ Important principle, performed at many levels in a computer (in hardware, operating system, software)
- ❑ Information in use **copied from slower to faster storage temporarily**
- ❑ Faster storage (cache) checked first to determine if information is there
 - ❑ If it is, information used directly from the cache (fast)
 - ❑ If not, data copied to cache and used there
- ❑ Cache smaller than storage being cached
 - ❑ Cache management important design problem
 - ❑ Cache size and replacement policy





Books

- ❑ Operating Systems Concept
 - ❑ Written by Galvin and Silberschatz
 - ❑ Edition: 9th



References

- ❑ Operating Systems Concept
 - ❑ Written by Galvin and Silberschatz
 - ❑ Edition: 9th